

CADDI

Competitive comparison boundaries after the August product launch

A public-evidence decision brief prepared for Alejandro Castellano and the Caddi team.

Scope lock

Public sources only
No internal roadmap or win/loss data
Competitor pages checked against primary competitor sources
Findings separated as FACT / INFERENCE / UNKNOWN in Appendix

COMPLIMENTARY GIFT

Executive Brief + Evidence Appendix

Why we looked

Caddi launched a broader agent-building product while several public comparison pages still describe the May competitive boundary.



What changed

The review did not look for “new competitors.” It tested whether Caddi’s public comparison layer still draws the right boundaries.

Why it matters

Comparison pages are sales assets. If the boundary has moved, the durable differentiation should be re-tested before buyers do it themselves.

Executive finding

Three public comparison pages draw competitive boundaries narrower than the underlying evidence supports.

The issue is not that Caddi lacks differentiation.

It is that three public comparisons lean on boundaries that have become less clean: one competitor capability predated the page, and two competitor moves arrived after publication.

INFERENCE: Caddi may benefit from re-testing which differentiators are durable enough to remain central in its comparison story after the August launch.

Gumloop

Pre-existing gap

Gummie build/update capability predates Caddi page

Zapier

Post-publication drift

AI by Zapier moved agentic execution into the Zap editor

UiPath

Post-publication drift

Delegate adds describe/say/demonstrate under platform controls

Finding 1 - Gumloop

Caddi's page frames Gumloop as a build-and-maintain-yourself canvas, but Gummie already reduced both build and update burden.

CADDI PAGE - 29 MAY 2026

Gumloop is framed as a “DIY AI canvas” where the user wires nodes, builds the workflow, and owns maintenance.

“You build it” / “You maintain it” framing
Caddi contrast: “It’s built for you”
Source page still points to 70+ tools and May-era language

FACT

GUMLOOP SOURCE - 16 SEP 2025

Gummie lets a user describe a workflow and have Gumloop plan it, create it, connect nodes/variables, write prompts, and update existing workflows.

Build-by-conversation predates Caddi’s comparison page
The maintenance claim is also narrower: Gummie supports conversational updates
Caddi may still differ on vertical focus, deterministic production, and vendor-owned service

FACT

INFERENCE

The remaining distinction is narrower than “you build it there vs. Caddi builds it.”

Finding 2 - Zapier

The Zapier comparison was published before AI by Zapier moved agentic execution into the core Zap editor.

CADDI - 29 MAY

Zapier is described as simple app-to-app trigger/action glue that the customer builds and owns.

FACT

[C3]

ZAPIER - 15 JULY

AI by Zapier brings agentic tool calling, AI reasoning and autonomous task execution directly into a Zap step.

FACT

[Z1]

BOUNDARY SHIFT

Agentic + AI + deterministic steps can sit in one Zap
Per-tool approval exists for sensitive actions
Zap History records tools, data, model tier and task usage

POST-PUBLICATION DRIFT

INFERENCE



Finding 3 - UiPath

Delegate is a post-publication shift that narrows part of Caddi’s teach-without-building distinction.

CADDI PAGE - 29 MAY 2026

UiPath framed as classic RPA: developer-built bots
replaying screen clicks
Caddi framed as record-to-code automation that
runs over APIs
Caddi best-fit: non-technical ops teams in law and
finance

C4 FACT

UIPATH DELEGATE - 25 AUG 2026

**Delegate is in Public Preview and can take task
input through natural language, voice or screen
recording while operating under UiPath identity,
permissions and audit controls.**

The runtime remains different: UiPath calls it a desktop AI agent / native
computer use
This does not establish technical equivalence with Caddi
It does narrow part of the “teach it without building it” contrast

FACT INFERENCE

PUBLIC PREVIEW NOT EQUIVALENT

What this does and does not say

The findings are intentionally narrower than a full positioning critique.

Supported by evidence

One Caddi comparison under-describes a competitor capability that already existed when the page was published.
Two other comparisons have been overtaken by later, dated competitor product moves.
Caddi has real remaining differentiators that should be separated from claims that have become easier to qualify.

Not established

That Caddi lacks differentiation.
That Gumloop, Zapier or UiPath are equivalent to Caddi.
That Caddi's August launch is strategically wrong.
That buyer trust has already been damaged.
That internal roadmap or win/loss data would support the same view.

Decision area

The useful question is which differentiators should remain the center of the comparison story.

Undocumented work discovery

Can Caddi see and rank repeatable work before the customer writes a spec?

Exception extraction

Can the product extract edge cases while the process owner is present?

Deterministic execution

Where APIs exist, can Caddi deliver lower runtime variability than agent-only approaches?

Vendor-owned maintenance

Which tier / use case truly receives Caddi-owned upkeep rather than customer ownership?

Regulated workflow depth

Which proof belongs to law, wealth and insurance instead of generic automation?

INFERENCE

The next decision is not “which page needs edits.” It is which claims are durable enough to carry the post-launch story after a technical buyer checks the evidence.

Suggested actions

A small comparison-layer reset could protect the stronger claims Caddi already appears to have.

1

Rewrite the Gumloop comparison around the narrower distinction

Do not frame Gumloop as only manual canvas-building. Separate conversational build/update from Caddi's vertical, governed, maintained outcome.

2

Update Zapier after AI by Zapier GA

Treat Zapier as moving beyond simple glue inside the core Zap editor. Re-center on unstructured inputs, regulated work and ownership model.

3

Qualify the UiPath contrast

Keep the runtime difference, but account for Delegate's demonstration input and governance controls. Label Public Preview clearly.

4

Build a durable claim map

For each comparison claim: keep, qualify, retire, or monitor. Tie each to source evidence and a buyer question.

The next useful test is not a broader competitor list.

It is a decision sprint on which competitive claims should survive, change, or be watched after the August product launch.

If useful, reply with the comparison or decision you want tested next.

Evidence Axis can convert this into a Competitive Decision Sprint: claim map, comparison rewrite logic, source-backed battlecard and monitoring list.