

Keenable

Evidence Appendix

Pricing transition, performance proof, and workload-specific competitive pressure

Scope

Public evidence only. This appendix supports the accompanying Executive Intelligence Deck. It does not use internal Keenable data, customer contracts, private benchmark results, or win/loss evidence. Public pricing comparisons are list-price analysis, not workload-equivalence claims.

Prepared by Evidence Axis

Checked: 29 August 2026

What this appendix resolves

- The evidence chain behind Finding A: the free-to-production pricing handoff.
- The evidence chain behind Finding B: why independent evidence supports latency more clearly than a universal quality-leadership claim.
- Competitive-set validation for Parallel, Exa, and Firecrawl.
- Source classes, review dates, assumptions, and unresolved unknowns.

1. Evidence map

The executive deck contains two decision-relevant findings. Each is separated into fact, Evidence Axis analysis/inference, and unknowns that public evidence cannot resolve.

Finding	Element	Class	Evidence / interpretation
A	Free-to-production pricing handoff	FACT	Keenable offers 100K free requests/month; Agent Builder is \$4/1K; Frontier is \$1/1K at 100 RPS+. Parallel Fast is \$1/1K after its recurring free allowance.
A	Public-price crossover	ANALYSIS	Using recurring public list prices, Keenable and Parallel Fast reach equal monthly spend at about 131,667 requests/month.
A	Commercial exposure	INFERENCE	The likely pressure point is not trial acquisition. It is the transition from a successful free builder into recurring production usage.
B	Broad outside benchmark	FACT	Artificial Analysis launch data scores Keenable Realtime 67 versus Parallel Advanced 75, Exa Auto 74, and Firecrawl 73; Keenable has the fastest individual search calls at 0.34s.
B	Benchmark provenance boundary	FACT	NEEDLE is company-operated and methodologically transparent; Artificial Analysis is an independent fixed-harness benchmark. Their scores are not directly interchangeable.
B	GTM proof choice	INFERENCE	The cleanest broad outside proof currently centers on per-query latency; quality claims are safer when bounded to a named benchmark or workload.

Decision boundary

Public evidence can identify where the decision sits. It cannot tell us Keenable's actual conversion curve, customer willingness to pay, Frontier qualification rate, internal benchmark results, or which buyer segment founders intend to prioritize.

Source handling

Vendor primary sources establish product, pricing, hiring, and company-stated goals. Independent or third-party benchmark sources are used for external performance evidence. Company-run NEEDLE is treated as primary benchmark evidence with its self-operation disclosure preserved.

2. Finding A - free-to-production pricing handoff

FACT CHAIN

Source	Verified public evidence
S2 - Keenable pricing	100K free requests per month. Agent Builder: \$4 / 1,000, pay as you go, cloud-only. Frontier: \$1 / 1,000 at 100 RPS+, dedicated capacity, cloud and on-premises, positioned for AI labs and inference platforms.
S3 - Head of Product	The first dedicated product hire is asked to design the PLG wedge that makes Keenable the default for agent builders, differentiated on latency, cost and scale, and to chart a path to 10M+ diverse daily agent queries.
S4 - Head of Revenue	The first dedicated revenue hire is asked to shape pricing and packaging with the founders. A sample problem spans a solo agent builder doing hundreds of queries a day to a lab doing millions.
S7 - Parallel pricing	Parallel currently states up to 5,000 requests per month free. Search Fast and Turbo are \$1 / 1,000 requests with 10 results; Fast is positioned for interactive agents under a one-second budget.

What is non-obvious here

The pricing pages alone are obvious. The decision-relevant finding comes from connecting those prices to Keenable's public PLG and traffic-learning goals, then quantifying where the recurring public cost relationship changes.

Finding A

Keenable removes most early trial friction with a large recurring free allowance. On current recurring public list prices, the cost advantage versus Parallel Fast flips shortly after a builder moves beyond the free tier. That makes the production handoff - not initial activation - the part worth testing with internal usage data.

What this does not claim

It does not claim Keenable should cut price, that customers see the APIs as equivalent, that Parallel wins total task cost, or that the \$1 Frontier tier is unavailable to a specific customer. Those require internal or workload-level evidence.

3. Finding A - public list-price model

Model basis: recurring public allowances and public Search API list prices checked 29 Aug 2026. Q is monthly search requests. Signup promotions and negotiated contracts are excluded.

Cost functions

Keenable Agent Builder: $\max(0, Q - 100,000) / 1,000 \times \4 .

Parallel Fast: $\max(0, Q - 5,000) / 1,000 \times \1 .

For $Q > 100,000$, the public-price equality occurs at approximately 131,667 requests/month.

Monthly requests	Keenable Agent Builder	Parallel Fast	Difference
100K	\$0	\$95	Keenable lower by \$95
132K	\$128	\$127	Approximately equal
250K	\$600	\$245	Keenable higher by \$355
500K	\$1,600	\$495	Keenable higher by \$1,105
1M	\$3,600	\$995	Keenable higher by \$2,605

Interpretation guardrail

Allowed: “On recurring public list prices, the Keenable Agent Builder / Parallel Fast crossover is about 132K monthly requests.”

Not allowed: “Parallel is four times cheaper in production” or “Keenable pricing blocks adoption.” The two products can differ in result shape, workload behavior, task-level call counts, contracts, and eligibility for higher-volume tiers.

Unknowns that matter before action

- Distribution of Keenable users across monthly request bands.
- Share of growing builders that move to Frontier or negotiated terms.
- Price elasticity and churn/upgrade behavior after free usage.
- Comparable downstream cost per completed agent task, not just per search request.

4. Finding B - external proof pattern

FACT CHAIN

Surface	What it measures	Verified result relevant to Keenable
S5 - NEEDLE	Live, company-operated search benchmark across several verticals; same query text; blind shared judge; public raw artifacts.	Method is unusually transparent. Keenable discloses that it operates NEEDLE and has not yet performed a systematic human-agreement audit of the judge.
S6 - Artificial Analysis	Fixed agent/model harness; search provider swapped. Broad QA/research tasks.	Launch: Keenable Realtime 67; Parallel Advanced 75; Exa Auto 74; Firecrawl 73. Current AA FAQ identifies Keenable Search (realtime) at 0.34s average per individual search query, fastest in its benchmark. Live-board rankings can change.

Finding B | Independent public evidence supports Keenable's speed advantage more clearly than broad quality leadership. Artificial Analysis does not rank Keenable first on its broad Search Index, while it independently verifies the fastest individual search calls at 0.34s. NEEDLE supports a quality case under a separate company-operated live methodology. The safer GTM architecture is therefore latency/scale as broad proof, with quality claims explicitly bounded to a named benchmark or workload.

5. Benchmark interpretation - what can and cannot be compared

The following measurements should not be merged into one composite ranking. They answer different questions.

Benchmark	Unit / metric	Useful for	Do not infer
Artificial Analysis	Search Index; per-query search latency; task-level agent metrics	Independent broad agent performance under one fixed model/harness	That a 67 score means Keenable is weak on every retrieval workload, or that live rankings cannot change.
NEEDLE	Vertical-specific ranking/recall metrics on fresh query streams	Freshness, live query behavior and company-run benchmark-specific quality proof	That a company-operated benchmark is independent third-party validation, or directly comparable to AA scores.

Why this matters commercially

A buyer can encounter different “best” providers depending on benchmark design and workload. Keenable can reduce claim risk by separating: (1) independently validated latency/scale proof, (2) workload-specific quality proof, and (3) company-run live benchmark evidence.

Claim architecture suggested by the evidence

Broad claim: low-latency agent-native search at scale.

Bounded quality claim: state the named benchmark, vertical, or workload.

Avoid: a universal “best quality” reading unless independent evidence converges across the workloads Keenable wants to own.

6. Competitive set and workload movement

Parallel, Exa, and Firecrawl remain the focused competitor set because independent benchmark surfaces place them in the same search-API buying category, while their recent product moves sharpen distinct agent jobs.

Competitor	Independent competitive-set support	Recent primary-source move	Decision relevance
Parallel	Artificial Analysis includes Parallel and Keenable in one Search API benchmark. Openbenchmarks separately measures Parallel in a fixed research-agent search workflow alongside Exa and Firecrawl.	Fast mode launched 21 Aug at \$1/1K. Current pricing also lists Turbo at \$1/1K and recurring free usage.	Price/speed pressure for general agent search.
Exa	Artificial Analysis includes Exa and Keenable. Openbenchmarks separately measures Exa with Parallel and Firecrawl.	23 Jul: dedicated search over roughly 350M publications.	Specialized research-retrieval wedge and proof surface.
Firecrawl	Artificial Analysis includes Firecrawl and Keenable. Openbenchmarks separately measures Firecrawl with Parallel and Exa.	20 Aug: Developer Index with 70M+ developer artifacts.	Specialized coding-agent retrieval wedge.

Keenable distribution signal

On 28 Aug, Keenable announced that Lightpanda Agent now uses Keenable when no search key is configured, with the client contributed upstream. This is evidence of distribution work that can place Keenable inside an agent workflow before an explicit vendor-selection step.

Implication

Competitors are not relying on a single “search API” story. They are claiming specific jobs. Keenable can keep the broader infrastructure ambition while choosing one easy-to-verify general wedge and building specialized proof where it wants to counter those claims.

7. Fact / inference / unknown register

ID	Class	Statement
R1	FACT	Keenable pricing currently gives 100K free requests/month, then Agent Builder at \$4/1K; Frontier is \$1/1K at 100 RPS+.
R2	ANALYSIS	Recurring public list-price equality with Parallel Fast occurs at ~131.7K requests/month under the stated assumptions.
R3	INFERENCE	The pricing risk, if real, is more likely after activation than at trial entry.
R4	UNKNOWN	Actual customer usage distribution, negotiated terms, Frontier routing, price sensitivity, and retention are not public.
R5	FACT	Artificial Analysis independently verifies Keenable Realtime as fastest per individual search query at 0.34s in its current benchmark FAQ.
R6	FACT	Artificial Analysis does not rank Keenable first on its broad Search Index; launch score for Realtime is 67.
R7	FACT	NEEDLE is company-operated; Keenable discloses that it has not yet completed a systematic human-agreement audit of the judge. NEEDLE and Artificial Analysis use different methods and their scores are not directly comparable.
R8	INFERENCE	The most defensible broad GTM proof line currently centers on latency/scale, while quality claims should be benchmark-bounded.
R9	UNKNOWN	Public evidence does not establish which performance axis Keenable buyers value most or which internal benchmarks founders use for positioning.

Contradictions / scope notes

- Artificial Analysis launch and live-board values are time-stamped snapshots; the live board can change.
- NEEDLE and Artificial Analysis use different tasks and methods. Their scores are not directly comparable.
- NEEDLE is company-operated and should not be presented as independent third-party validation.
- Public price-per-request comparisons do not equal total cost per completed agent task.

8. Source register

Review date for all sources: 29 Aug 2026. URLs below are the pages reviewed for this package.

S1 Independent press	TechCrunch - Accel-backed Keenable is indexing the web for AI agents https://techcrunch.com/2026/08/25/accel-backed-keenable-is-indexing-the-web-for-ai-agents/ Used for: Seed round, team size, planned headcount/GTM expansion.
S2 Primary / Keenable	Keenable Pricing https://keenable.ai/pricing Used for: Free allowance, Agent Builder and Frontier public pricing/segment language.
S3 Primary / Keenable	Head of Product - Keenable https://app.keenable.ai/careers/head-of-product Used for: PLG wedge, latency/cost/scale differentiation, 10M+ diverse daily query goal.
S4 Primary / Keenable	Head of Revenue - Keenable https://app.keenable.ai/careers/head-of-revenue Used for: Pricing/package responsibility and solo-builder-to-lab sample problem.
S5 Primary benchmark / Keenable	NEEDLE benchmark and methodology https://keenableai.github.io/needle/ Used for: Live methodology, engine parity rules, self-operation disclosure and judge-audit limitation.
S6 Independent benchmark	Artificial Analysis Search Index - launch article and live board https://artificialanalysis.ai/articles/search-api Used for: Launch scores, task/latency explanation, Keenable launch measurements. Live board also checked: https://artificialanalysis.ai/agents/search-api
S7 Primary / competitor	Parallel pricing and Fast launch https://parallel.ai/pricing Used for: Recurring free allowance and current Search API tiers. Fast launch also checked: https://parallel.ai/blog/parallel-search-fast
S8 Independent benchmark platform	Openbenchmarks - Multi Turn Company Search https://openbenchmarks.com/multi-turn-company-search Used for: Independent competitive-set support; fixed-agent comparison includes Parallel, Exa and Firecrawl across 45 questions and repeated runs. Public rankings state that vendors do not pay for inclusion, ranking, or removal.

S9
Primary / competitor

Exa - SOTA Search Over Academic Publications

<https://exa.ai/blog/publications-search>
Used for: Dedicated index of roughly 350M publications; launched 23 Jul 2026.

S10
Primary / competitor

Firecrawl Developer Index launch

<https://www.firecrawl.dev/blog/developer-index-launch>
Used for: 70M+ developer artifacts; launched 20 Aug 2026.

S11
Primary / Keenable

Keenable now runs inside Lightpanda

<https://keenable.ai/blog/keenable-now-runs-inside-lightpanda>
Used for: Lightpanda Agent keyless-fallback integration and upstream client contribution.

Evidence Axis separates facts from analysis and unknowns. No public source in this appendix establishes Keenable's internal intent, private customer economics, or private win/loss evidence.