

Keenable

GTM timing is good. The handoff into production matters now.

A public-evidence brief on the free-to-production pricing handoff, externally validated performance, and workload-specific competitive pressure after Keenable's stealth exit.

READ THIS AS

A decision brief, not a company summary.

\$26M

Seed round reported at stealth exit


2

Decision exposures that survived the attack step

Prepared by Evidence Axis

Why we looked now

The timing is not just funding. Keenable is turning public evidence into GTM proof while its market frame is still being set.



Decision exposure Keenable is still choosing how to package price, quality, latency and segment focus. That makes public competitive evidence useful before the motion hardens.

Executive readout

Two findings survived the attack step. The first is quantified; the second survives without relying on unverified third-party reproductions.

FINDING A

The public cost advantage flips near 132K requests/month

Keenable's 100K free monthly allowance wins early trial economics. On current recurring public list prices, a growing builder crosses into higher spend versus Parallel Fast at roughly 132K monthly requests - before Frontier's \$1/1K high-throughput tier becomes the obvious path.

ANALYSIS + INFERENCE

FINDING B

Independent evidence supports a speed advantage more clearly than broad quality leadership

Artificial Analysis does not show a broad quality lead for Keenable on its fixed harness, but it does independently verify the fastest individual search calls at 0.34s. NEEDLE supports a quality case under a separate, company-operated live methodology.

FACT + INFERENCE

Recommended direction: keep the generous free trial, examine the production-price handoff, and lead GTM with latency plus workload-bounded quality proof rather than one broad SOTA claim.

Finding A - the free-to-production cost advantage flips early

The 100K free allowance is a real acquisition advantage. The pressure point appears when successful builders graduate into recurring production usage.

PUBLIC RECURRING PRICE PATH

Keenable trial

100K requests / month free
Then Agent Builder: \$4 / 1K

Keenable Frontier

\$1 / 1K at 100 RPS+
Dedicated capacity; AI
labs / inference

Parallel Search Fast

5K requests / month free
Then \$1 / 1K

**The question is not trial friction.
It is the price handoff after a builder
succeeds.**

PUBLIC LIST-PRICE MODEL

100K requests/mo

\$0 Keenable
\$95 Parallel Fast

250K requests/mo

\$600 Keenable
\$245 Parallel Fast

500K requests/mo

\$1,600 Keenable
\$495 Parallel Fast

1M requests/mo

\$3,600 Keenable
\$995 Parallel Fast

**~132K
requests/
mo
public-price
crossover**

Analysis uses recurring public list prices checked 29 Aug 2026: Keenable 100K free then \$4/1K; Parallel 5K free then Fast \$1/1K. It excludes signup promotions, custom contracts, Frontier eligibility, and differences in result shape or downstream task cost.

Decision question: should the production handoff be redesigned before builder usage becomes valuable feedback at scale?

Why Finding A matters to Keenable specifically

This is not a pricing-page nit. It intersects with the exact PLG and pricing problems Keenable is hiring its first product and revenue leaders to solve.

COMPANY-STATED PROBLEM


Keenable wants the product hooks and feedback loops that move from limited data to 10M+ diverse daily agent queries.

The same role asks for a PLG wedge that makes Keenable the default for agent builders, differentiated on latency, cost and scale.

Source: Head of Product role

TENSION

The tension starts after activation. Keenable wants more diverse production traffic and explicit feedback to improve retrieval.



At public list prices, the transition becomes more expensive than Parallel Fast at ~132K monthly requests. Whether that matters depends on real usage distribution and conversion behavior.

REVENUE ROLE MATCH

The first revenue hire is explicitly tasked with shaping pricing and packaging for usage-based search and fetch APIs with the founders.

One sample problem: price a usage-based search API spanning a solo builder doing hundreds of queries a day to a lab doing millions.

This makes the pricing handoff a live decision, not a retrospective critique.

Finding B - independent evidence supports a speed-first proof line

NEEDLE remains useful and transparent. The decision issue is claim architecture: company-run evidence supports a quality story, while the independent AA harness validates low per-query latency more cleanly than broad quality leadership.

THREE VERIFIED PROOF ELEMENTS

COMPANY-RUN

NEEDLE - live quality benchmark

Live query streams, blind shared judge and public raw artifacts.

Keenable discloses that it operates NEEDLE and has not yet completed a systematic human-agreement audit of the judge.

Useful benchmark-specific proof; not independent third-party validation.

INDEPENDENT

Artificial Analysis - quality

Launch Search Index: Keenable Realtime 67 vs Parallel Advanced 75, Exa Auto 74 and Firecrawl 73.

AA swaps the search provider inside a fixed agent/model harness.

On this harness, Keenable does not lead broad aggregate quality.

INDEPENDENT

Artificial Analysis - latency

Current AA FAQ: Keenable Search (realtime) averages 0.34s per individual search query - fastest in the benchmark.

This is the cleanest externally validated performance claim in the public evidence.

Live-board rankings can change; this claim is time-stamped to the checked source.

Inference: lead broad GTM proof with latency/scale, then state quality claims against a named benchmark or workload. NEEDLE and AA use different methods and should not be merged into one universal quality score.

Competitors are narrowing by workload

The pressure is not only benchmark rank. Nearby providers are making sharper claims around specific agent jobs.

Parallel

Cheap, fast general agent search

Fast mode launched Aug 21 at \$1/1K.
Latest Artificial Analysis (Aug 27) ranks
Fast fastest per task and Advanced
highest on the Search Index.

S6, S7

Exa

Research-heavy retrieval

Dedicated publication search over
roughly 350M papers and technical
publications gives Exa a visible use-
case claim for research agents.

S9

Firecrawl

Developer artifacts for coding agents

Developer Index covers 70M+
READMEs, issues, pull requests, docs
and OpenAPI specs, with result
passages available to agents.

S10

Implication for Keenable: competitors are owning specific jobs. A broad infrastructure story needs one general proof wedge plus workload-specific proof where Keenable wants to win.

Decision matrix

The question is not whether Keenable can win. It is which claims should carry the first buyer conversation while the proof stack is still being built.

Option	Evidence strength	Risk	Best use
Lead with broad SOTA quality	NEEDLE supports a quality case under its own live methodology; AA's independent fixed harness has other leaders on its Search Index.	Universal wording can be challenged with independent AA evidence; NEEDLE is company-operated and should be labeled as such.	Use only with benchmark-specific wording.
Lead with latency + scale	AA independently verifies 0.34s average per individual search query - fastest in its benchmark.	Narrower than the full product ambition, but difficult to dispute.	Best general near-term proof wedge.
Build workload-specific quality proof	NEEDLE and AA use different tasks and methods; their quality scores are not interchangeable.	Requires separate external evals for the target jobs Keenable wants to own; less convenient than one SOTA headline.	Best proof architecture alongside latency.

Evidence Axis recommendation: lead with independently verified latency/scale now; keep quality claims benchmark-bounded; build external workload-specific proof before using one universal SOTA reading.

Recommended next checks

Three checks would turn the public evidence into decisions using Keenable's internal usage and buyer data.

1

Production-price handoff

Map actual monthly usage distribution around the ~132K public-price crossover. Test retention, upgrade and Frontier routing before changing list price.

2

Benchmark claim architecture

Separate broad retrieval, coding, research and other priority workloads. Define which claim is supported by third-party evidence for each one.

3

Competitive workload response

Map the jobs Parallel, Exa and Firecrawl are claiming, then choose where Keenable should counter, ignore, or create a stronger specialist proof surface.

If useful, reply with the next competitor or decision you want tested.

Source key

Key sources used in the executive deck. The accompanying Evidence Appendix contains the full register, URLs, source class, review date, and caveats.

S1	TechCrunch, 25 Aug: \$26M seed; 15 engineering staff; planned GTM headcount expansion.	S8	Openbenchmarks company-search board: independently compares Parallel, Exa and Firecrawl in one fixed research-agent workflow.
S2	Keenable pricing, checked 29 Aug: 100K free/mo; Agent Builder \$4/1K; Frontier \$1/1K at 100 RPS+.	S9	Exa Publications, 23 Jul: dedicated index of ~350M publications.
S3	Keenable Head of Product: PLG wedge; latency/cost/scale; path to 10M+ diverse daily queries.	S10	Firecrawl Developer Index, 20 Aug: 70M+ developer artifacts.
S4	Keenable Head of Revenue: shape pricing/packaging; solo builder to lab doing millions.	S11	Keenable Lightpanda integration, 28 Aug: keyless fallback search integration shipped upstream.
S5	NEEDLE: live/open benchmark; blind judge; raw artifacts; self-operated; no systematic human-agreement audit yet.		
S6	Artificial Analysis Search Index launch + live board (latest data Aug 27): quality, cost and latency under a fixed agent harness.		
S7	Parallel pricing + Fast launch: recurring 5K free/mo; Fast \$1/1K; launch 21 Aug.		

Method note: cost crossover is Evidence Axis analysis from recurring public list prices. Benchmark metrics are not treated as interchangeable across different test designs.